

▼ Part I - Exploring the Prosper Loan Dataset

Haneen Habash

My Github: [@haneenhabash](#)

My LinkedIn: [@haneen-habash-26b226282](#)

▼ Table of Contents

- Introduction
 - Data Dictionary
 - Objectives
 - Premirely Wrangling
 - Univariate extrapolation
 - Bivariate extrapolation
 - Multivariate extrapolation
 - Conclusions
-

▼ Introduction

On November 24, 2008, the SEC found Prosper to be in violation of the Securities Act of 1933. As a result of these findings, the SEC imposed a cease and desist order on Prosper ... In July 2009, Prosper reopened their website for lending ("investing") and borrowing after having obtained SEC registration for its loans ("notes"). After the relaunch, bidding on loans was restricted to residents of 28 U.S. states and the District of Columbia. Borrowers may reside in any of 47 states, with residents of three states (Iowa, Maine, and North Dakota) not permitted to borrow through Prosper

In this notebook, the analysis is done on the Prosper Datatset which is collected from a Loan company. The dataset includes customers who have paid off their

loans, who have been past due and put into collection without paying back their loan and interests, and who have paid off only after they were put in collection. The original dataset contains 113937 rows and 81 columns out of which 12 features of interest were selected.

▼ Data Dictionary

The following data dictionary shows each variable of the dataset and the corresponding description:

Variable	
ListingKey	Unique key for each listing, same value as the 'key' used in the listing object in the API.
ListingNumber	The number that uniquely identifies the listing to the public as displayed on the website.
ListingCreationDate	The date the listing was created.
CreditGrade	The Credit rating that was assigned at the time the listing went live. Applicable for listing.
Term	The length of the loan expressed in months.
LoanStatus	The current status of the loan: Cancelled, Chargedoff, Completed, Current, Defaulted, Failed.
ClosedDate	Closed date is applicable for Cancelled, Completed, Chargedoff and Defaulted loan status.
BorrowerAPR	The Borrower's Annual Percentage Rate (APR) for the loan.
BorrowerRate	The Borrower's interest rate for this loan.
LenderYield	The Lender yield on the loan. Lender yield is equal to the interest rate on the loan less the fees.
EstimatedEffectiveYield	Effective yield is equal to the borrower interest rate (i) minus the servicing fee rate, (ii) net of principal prepayments.
EstimatedLoss	Estimated loss is the estimated principal loss on charge-offs. Applicable for loans originated by Prosper.
EstimatedReturn	The estimated return assigned to the listing at the time it was created. Estimated return is the estimated yield plus the estimated loss.
ProsperRating (numeric)	The Prosper Rating assigned at the time the listing was created: 0 - N/A, 1 - HR, 2 - E, 3 - AA, 4 - A, 5 - B, 6 - C, 7 - D, 8 - F, 9 - G, 10 - I.
ProsperRating (Alpha)	The Prosper Rating assigned at the time the listing was created between AA - HR. Applied to loans originated by Prosper.
ProsperScore	A custom risk score built using historical Prosper data. The score ranges from 1-10, with 1 being the lowest risk and 10 being the highest risk.
ListingCategory	The category of the listing that the borrower selected when posting their listing: 0 - Not Specified, 1 - Personal, 2 - Business, 3 - Investment, 4 - Home Improvement, 5 - Other.
BorrowerState	The two letter abbreviation of the state of the address of the borrower at the time the listing was created.
Occupation	The Occupation selected by the Borrower at the time they created the listing.
EmploymentStatus	The employment status of the borrower at the time they posted the listing.
EmploymentStatusDuration	The length in months of the employment status at the time the listing was created.
IsBorrowerHomeowner	A Borrower will be classified as a homeowner if they have a mortgage on their credit profile.
CurrentlyInGroup	Specifies whether or not the Borrower was in a group at the time the listing was created.
GroupKey	The Key of the group in which the Borrower is a member of. Value will be null if the borrower is not in a group.
DateCreditPulled	The date the credit profile was pulled.
CreditScoreRangeLower	The lower value representing the range of the borrower's credit score as provided by a credit bureau.
CreditScoreRangeUpper	The upper value representing the range of the borrower's credit score as provided by a credit bureau.
FirstRecordedCreditLine	The date the first credit line was opened.
CurrentCreditLines	Number of current credit lines at the time the credit profile was pulled.

Variable	
OpenCreditLines	Number of open credit lines at the time the credit profile was pulled.
TotalCreditLinespast7years	Number of credit lines in the past seven years at the time the credit profile was pulled.
OpenRevolvingAccounts	Number of open revolving accounts at the time the credit profile was pulled.
OpenRevolvingMonthlyPayment	Monthly payment on revolving accounts at the time the credit profile was pulled.
InquiriesLast6Months	Number of inquiries in the past six months at the time the credit profile was pulled.
TotalInquiries	Total number of inquiries at the time the credit profile was pulled.
CurrentDelinquencies	Number of accounts delinquent at the time the credit profile was pulled.
AmountDelinquent	Dollars delinquent at the time the credit profile was pulled.
DelinquenciesLast7Years	Number of delinquencies in the past 7 years at the time the credit profile was pulled.
PublicRecordsLast10Years	Number of public records in the past 10 years at the time the credit profile was pulled.
PublicRecordsLast12Months	Number of public records in the past 12 months at the time the credit profile was pulled
RevolvingCreditBalance	Dollars of revolving credit at the time the credit profile was pulled.
BankcardUtilization	The percentage of available revolving credit that is utilized at the time the credit profile was pulled.
AvailableBankcardCredit	The total available credit via bank card at the time the credit profile was pulled.
TotalTrades	Number of trade lines ever opened at the time the credit profile was pulled.
TradesNeverDelinquent	Number of trades that have never been delinquent at the time the credit profile was pulled.
TradesOpenedLast6Months	Number of trades opened in the last 6 months at the time the credit profile was pulled.
DebtToIncomeRatio	The debt to income ratio of the borrower at the time the credit profile was pulled. This value will be null if the borrower does not have a verifiable income.
IncomeRange	The income range of the borrower at the time the listing was created.
IncomeVerifiable	The borrower indicated they have the required documentation to support their income.
StatedMonthlyIncome	The monthly income the borrower stated at the time the listing was created.
LoanKey	Unique key for each loan. This is the same key that is used in the API.
TotalProsperLoans	Number of Prosper loans the borrower at the time they created this listing. This value will be null if the borrower does not have a verifiable income.
TotalProsperPaymentsBilled	Number of on time payments the borrower made on Prosper loans at the time they created this listing.
OnTimeProsperPayments	Number of on time payments the borrower had made on Prosper loans at the time they created this listing.
ProsperPaymentsLessThanOneMonthLate	Number of payments the borrower made on Prosper loans that were less than one month late.
ProsperPaymentsOneMonthPlusLate	Number of payments the borrower made on Prosper loans that were greater than one month late.
ProsperPrincipalBorrowed	Total principal borrowed on Prosper loans at the time the listing was created. This value will be null if the borrower does not have a verifiable income.
ProsperPrincipalOutstanding	Principal outstanding on Prosper loans at the time the listing was created. This value will be null if the borrower does not have a verifiable income.
ScorexChangeAtTimeOfListing	Borrower's credit score change at the time the credit profile was pulled. This will be the difference between the current score and the score at the time the listing was created.
LoanCurrentDaysDelinquent	The number of days delinquent.
LoanFirstDefaultedCycleNumber	The cycle the loan was charged off. If the loan has not charged off the value will be null.
LoanMonthsSinceOrigination	Months since the loan originated.
LoanNumber	The number that uniquely identifies the loan to the public as displayed on the website.
LoanOriginalAmount	The original amount of the loan.
LoanOriginationDate	The date the loan originated.
LoanOriginationQuarter	The quarter in which the loan originated.
MemberKey	Unique key for each member. This is the same key that is used in the API.
MonthlyLoanPayment	The monthly payment (principal and interest) the borrower is required to make for this loan.

Variable	
LP_CustomerPayments	The total payments (principal + interest) that have been made on the loan by the borrower.
LP_CustomerPrincipalPayments	The total principal payments that have been made on the loan by the borrower.
LP_InterestandFees	Interest and fees paid by the borrower.
LP_ServiceFees	The servicing fees paid by the borrower.
LP_CollectionFees	The collection fees paid by the borrower.
LP_GrossPrincipalLoss	Gross principal loss on the loan.
LP_NetPrincipalLoss	Net principal loss on the loan.
LP_NonPrincipalRecoverypayments	Non-principal recovery payments on the loan.
PercentFunded	The percentage of the loan that was funded.
Recommendations	Number of recommendations the borrower had at the time they created the listing.
InvestmentFromFriendsCount	Number of investments that were made by friends at the time the listing was created.
InvestmentFromFriendsAmount	The dollar amount of investments that were made by friends at the time the listing was created.
Investors	The number of investors that funded the loan.

▼ Objectives

- 1. Loan Performance Analysis**
- 2. Credit Score and Borrower Analysis**
- 3. Geographic and Demographic Analysis**

▼ Preliminary Wrangling

- In this section, a preliminary data wrangling is done on the dataset.

```
## import all packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# suppress warnings from final output
import warnings
warnings.simplefilter("ignore")
```

```
→ c:\Users\SS\AppData\Local\Programs\Python\Python311\Lib\site-packages\scipy\__init__.py:
    warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}"
```

>Loading the dataset

Let's Load in the dataset into a pandas dataframe:

```
df = pd.read_csv("./data/prosperLoanData.csv") ## Load the csv into pandas dataframe
df.sample(10) ## Looking at a random sample of 10 rows.
```

89396	522435803124450551D167F	796219		2013-06-03 18:11:11.817000000		NaN	36
66329	CD2435638856954192AA8E9	682194		2012-12-04 06:24:02.113000000		NaN	36
57866	811B3553828282161FA98C0	616886		2012-07-25 20:22:43.937000000		NaN	36
68993	71C2340109095088293289C	204704		2007-09-21 00:36:36.917000000		AA	36
68782	FDB1357349001237607843C	729927		2013-03-14 16:51:11.260000000		NaN	36
79820	1CB13601952390901F76B99	1172289		2014-01-30 12:50:33.480000000		NaN	36
107870	2260360104242486260A5C5	1142878		2014-02-04 13:01:44.540000000		NaN	60
55104	BD9335910112976028271D0	969276		2013-10-09 19:02:28.670000000		NaN	60
106857	F4533393067041235BB5B50	157991		2007-06-25 16:35:41.160000000		AA	36
44207	5A1735791083505443D330B	785172		2013-05-21 06:09:04.963000000		NaN	60

10 rows × 81 columns

Dataset Structure

```
df.shape ## showing the shape of the dataset
```

→ (113937, 81)

This dataset has 113,937 rows and 81 columns. Which is a relatively big dataset.

▼ Dataset Assessment and Cleaning

▼ Duplicated Records

```
df.duplicated(subset='LoanKey').sum()
```

→ 871

Let's identify the duplicated records based on the LoanKey and see if we should handle this

```
duplicates = df[df.duplicated(subset='LoanKey', keep=False)]
display(duplicates.head(10))
```

		ListingKey	ListingNumber	ListingCreationDate	CreditGrade	Term	L
8	0F043596202561788EA13D5		1023355	2013-12-02 10:43:39.117000000	NaN	36	
9	0F043596202561788EA13D5		1023355	2013-12-02 10:43:39.117000000	NaN	36	
29	0F563597161095613517437		1051243	2013-12-17 09:18:33.220000000	NaN	36	
176	106335993636414276CB477		1119836	2014-01-08 14:27:50.320000000	NaN	36	
313	09233589620788733CFB8CE		930842	2013-09-25 08:03:11.860000000	NaN	36	
349	313635901230654318A9030		931467	2013-09-26 18:50:29.053000000	NaN	36	
442	09AD35918712001025AC1BD		969821	2013-10-24 13:21:31.607000000	NaN	36	
444	09CD3592594126374FB0A7C		986199	2013-10-18 08:28:03.610000000	NaN	36	
455	31C73597152310464749E00		1092437	2013-12-23 13:47:35.500000000	NaN	36	
461	44F2358557406858060EBDE		870200	2013-08-15 07:12:49.410000000	NaN	60	

10 rows × 81 columns

```
df.drop_duplicates(subset='LoanKey', keep='first', inplace=True) ## Dropping the duplicated
```

```
print(df.duplicated(subset='LoanKey').sum()) ## Checking the drop of duplicated
df.shape
```

→ 0
(113066, 81)

Let's check for the duplicated records based on the ListingKey based on the documentation it has to be unique too.

```
df.duplicated(subset='ListingKey').sum() ## checking for duliaited based on the listing key
```

→ 0

▼ Data types Validity

- Assessment: Let's look at the data types of these variables and assess them using .info:

```
df.info()
```

→ <class 'pandas.core.frame.DataFrame'>

Index: 113066 entries, 0 to 113936

Data columns (total 81 columns):

#	Column	Non-Null Count	Dtype	
0	ListingKey	113066	non-null	object
1	ListingNumber	113066	non-null	int64
2	ListingCreationDate	113066	non-null	object
3	CreditGrade	28953	non-null	object
4	Term	113066	non-null	int64
5	LoanStatus	113066	non-null	object
6	ClosedDate	55076	non-null	object
7	BorrowerAPR	113041	non-null	float64
8	BorrowerRate	113066	non-null	float64
9	LenderYield	113066	non-null	float64
10	EstimatedEffectiveYield	83982	non-null	float64
11	EstimatedLoss	83982	non-null	float64
12	EstimatedReturn	83982	non-null	float64
13	ProsperRating (numeric)	83982	non-null	float64
14	ProsperRating (Alpha)	83982	non-null	object
15	ProsperScore	83982	non-null	float64
16	ListingCategory (numeric)	113066	non-null	int64
17	BorrowerState	107551	non-null	object
18	Occupation	109537	non-null	object
19	EmploymentStatus	110811	non-null	object
20	EmploymentStatusDuration	105441	non-null	float64
21	IsBorrowerHomeowner	113066	non-null	bool
22	CurrentlyInGroup	113066	non-null	bool
23	GroupKey	13339	non-null	object
24	DateCreditPulled	113066	non-null	object
25	CreditScoreRangeLower	112475	non-null	float64
26	CreditScoreRangeUpper	112475	non-null	float64
27	FirstRecordedCreditLine	112369	non-null	object
28	CurrentCreditLines	105462	non-null	float64
29	OpenCreditLines	105462	non-null	float64
30	TotalCreditLinespast7years	112369	non-null	float64
31	OpenRevolvingAccounts	113066	non-null	int64
32	OpenRevolvingMonthlyPayment	113066	non-null	float64
33	InquiriesLast6Months	112369	non-null	float64
34	TotalInquiries	111907	non-null	float64
35	CurrentDelinquencies	112369	non-null	float64
36	AmountDelinquent	105444	non-null	float64
37	DelinquenciesLast7Years	112076	non-null	float64
38	PublicRecordsLast10Years	112369	non-null	float64
39	PublicRecordsLast12Months	105462	non-null	float64
40	RevolvingCreditBalance	105462	non-null	float64
41	BankcardUtilization	105462	non-null	float64

42	AvailableBankcardCredit	105522	non-null	float64
43	TotalTrades	105522	non-null	float64
44	TradesNeverDelinquent (percentage)	105522	non-null	float64
45	TradesOpenedLast6Months	105522	non-null	float64
46	DebtToIncomeRatio	104594	non-null	float64
47	IncomeRange	113066	non-null	object
48	IncomeVerifiable	113066	non-null	bool
49	StatedMonthlyIncome	113066	non-null	float64
50	LoanKey	113066	non-null	object
51	TotalProsperLoans	21923	non-null	float64
52	TotalProsperPaymentsBilled	21923	non-null	float64

We have the columns of `ClosedDate`, `LoanOriginationDate`, `DateCreditPulled`, and the `ListingCreationDate` has an object type and it has to be a datetime type.

```
## listing the date columns
date_columns = ['ClosedDate', 'LoanOriginationDate', 'DateCreditPulled', 'ListingCreationDat

#### Loopong through the list and covering to datetime data type
for col in date_columns:
    df[col] = pd.to_datetime(df[col], format='mixed') ## using the format as mixed to infer

df.info() #### check if that is successful
```

```
→ <class 'pandas.core.frame.DataFrame'>
Index: 113066 entries, 0 to 113936
Data columns (total 81 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ListingKey       113066 non-null   object 
 1   ListingNumber    113066 non-null   int64  
 2   ListingCreationDate  113066 non-null   datetime64[ns]
 3   CreditGrade      28953 non-null   object 
 4   Term             113066 non-null   int64  
 5   LoanStatus       113066 non-null   object 
 6   ClosedDate       55076 non-null   datetime64[ns]
 7   BorrowerAPR     113041 non-null   float64
 8   BorrowerRate     113066 non-null   float64
 9   LenderYield      113066 non-null   float64
 10  EstimatedEffectiveYield  83982 non-null   float64
 11  EstimatedLoss    83982 non-null   float64
 12  EstimatedReturn   83982 non-null   float64
 13  ProsperRating (numeric) 83982 non-null   float64
 14  ProsperRating (Alpha)  83982 non-null   object 
 15  ProsperScore      83982 non-null   float64
 16  ListingCategory (numeric) 113066 non-null   int64  
 17  BorrowerState     107551 non-null   object 
 18  Occupation        109537 non-null   object 
 19  EmploymentStatus   110811 non-null   object 
 20  EmploymentStatusDuration 105441 non-null   float64
 21  IsBorrowerHomeowner 113066 non-null   bool
```

22	CurrentlyInGroup	113066	non-null	bool
23	GroupKey	13339	non-null	object
24	DateCreditPulled	113066	non-null	datetime64[ns]
25	CreditScoreRangeLower	112475	non-null	float64
26	CreditScoreRangeUpper	112475	non-null	float64
27	FirstRecordedCreditLine	112369	non-null	object
28	CurrentCreditLines	105462	non-null	float64
29	OpenCreditLines	105462	non-null	float64
30	TotalCreditLinespast7years	112369	non-null	float64
31	OpenRevolvingAccounts	113066	non-null	int64
32	OpenRevolvingMonthlyPayment	113066	non-null	float64
33	InquiriesLast6Months	112369	non-null	float64
34	TotalInquiries	111907	non-null	float64
35	CurrentDelinquencies	112369	non-null	float64
36	AmountDelinquent	105444	non-null	float64
37	DelinquenciesLast7Years	112076	non-null	float64
38	PublicRecordsLast10Years	112369	non-null	float64
39	PublicRecordsLast12Months	105462	non-null	float64
40	RevolvingCreditBalance	105462	non-null	float64
41	BankcardUtilization	105462	non-null	float64
42	AvailableBankcardCredit	105522	non-null	float64
43	TotalTrades	105522	non-null	float64
44	TradesNeverDelinquent (percentage)	105522	non-null	float64
45	TradesOpenedLast6Months	105522	non-null	float64
46	DebtToIncomeRatio	104594	non-null	float64
47	IncomeRange	113066	non-null	object
48	IncomeVerifiable	113066	non-null	bool
49	StatedMonthlyIncome	113066	non-null	float64
50	LoanKey	113066	non-null	object
51	TotalProsperLoans	21923	non-null	float64

Now the rest of the data types of the variables are valid.

▼ Data Completeness

```
def get_percent_null(df):
    """
    Args:
        - df (pd.DataFrame)
    Returns:
        - The percentages of missing values
    """
    null_counts = df.isnull().sum()
    null_counts = null_counts[null_counts > 0].sort_values()
    return (null_counts/df.shape[0])*100
```

get_percent_null(df)

→ BorrowerAPR	0.022111
CreditScoreRangeUpper	0.522704

CreditScoreRangeLower	0.522704
PublicRecordsLast10Years	0.616454
CurrentDelinquencies	0.616454
InquiriesLast6Months	0.616454
TotalCreditLinespast7years	0.616454
FirstRecordedCreditLine	0.616454
DelinquenciesLast7Years	0.875595
TotalInquiries	1.025065
EmploymentStatus	1.994410
Occupation	3.121186
BorrowerState	4.877682
AvailableBankcardCredit	6.672209
TradesOpenedLast6Months	6.672209
TradesNeverDelinquent (percentage)	6.672209
TotalTrades	6.672209
CurrentCreditLines	6.725276
OpenCreditLines	6.725276
PublicRecordsLast12Months	6.725276
RevolvingCreditBalance	6.725276
BankcardUtilization	6.725276
AmountDelinquent	6.741195
EmploymentStatusDuration	6.743849
DebtToIncomeRatio	7.492969
ProsperRating (Alpha)	25.723029
ProsperRating (numeric)	25.723029
ProsperScore	25.723029
EstimatedLoss	25.723029
EstimatedEffectiveYield	25.723029
EstimatedReturn	25.723029
ClosedDate	51.288628
CreditGrade	74.392833
TotalProsperLoans	80.610440
TotalProsperPaymentsBilled	80.610440
OnTimeProsperPayments	80.610440
ProsperPaymentsLessThanOneMonthLate	80.610440
ProsperPaymentsOneMonthPlusLate	80.610440
ProsperPrincipalBorrowed	80.610440
ProsperPrincipalOutstanding	80.610440
ScorexChangeAtTimeOfListing	83.273486
LoanFirstDefaultedCycleNumber	85.006987
GroupKey	88.202466

dtype: float64

There are many columns that has null values, some of which have a very high count of null values. The following columns were the highest percentage of null such as higher than 50%. These columns are dropped.

```
high_missing_percent = ['CreditGrade', 'ScorexChangeAtTimeOfListing', 'LoanFirstDefaultedCyc
'TotalProsperPaymentsBilled', 'OnTimeProsperPayments', 'ProsperPaymentsLessThanOneMonthLate
'ProsperPrincipalBorrowed', 'ProsperPrincipalOutstanding'] ## list of columns with missing
```

```
df.drop(columns= high_missing_percent, inplace=True) ## Drop the list of high percentage col
```

```
df.shape ## Check if the dropping was successful
```

```
→ (113066, 70)
```

The drop is checked now that the variables went from 81 to 70.

Handling the other null values by dropping NA values and rows:

```
get_percent_null(df)
```

```
→ BorrowerAPR          0.022111  
CreditScoreRangeUpper 0.522704  
CreditScoreRangeLower 0.522704  
TotalCreditLinespast7years 0.616454  
CurrentDelinquencies 0.616454  
FirstRecordedCreditLine 0.616454  
PublicRecordsLast10Years 0.616454  
InquiriesLast6Months 0.616454  
DelinquenciesLast7Years 0.875595  
TotalInquiries        1.025065  
EmploymentStatus      1.994410  
Occupation            3.121186  
BorrowerState         4.877682  
AvailableBankcardCredit 6.672209  
TradesOpenedLast6Months 6.672209  
TotalTrades           6.672209  
TradesNeverDelinquent (percentage) 6.672209  
PublicRecordsLast12Months 6.725276  
RevolvingCreditBalance 6.725276  
BankcardUtilization   6.725276  
CurrentCreditLines    6.725276  
OpenCreditLines        6.725276  
AmountDelinquent      6.741195  
EmploymentStatusDuration 6.743849  
DebtToIncomeRatio     7.492969  
ProsperScore          25.723029  
ProsperRating (Alpha) 25.723029  
ProsperRating (numeric) 25.723029  
EstimatedReturn        25.723029  
EstimatedLoss          25.723029  
EstimatedEffectiveYield 25.723029  
GroupKey              88.202466  
dtype: float64
```

These columns are identifiers, mostly unique between each loan. They are irrelevant for the task and therefore they should be dropped too.

```
df.columns
```

```
→ Index(['ListingKey', 'ListingNumber', 'ListingCreationDate', 'Term',
       'LoanStatus', 'BorrowerAPR', 'BorrowerRate', 'LenderYield',
       'EstimatedEffectiveYield', 'EstimatedLoss', 'EstimatedReturn',
       'ProsperRating (numeric)', 'ProsperRating (Alpha)', 'ProsperScore',
       'ListingCategory (numeric)', 'BorrowerState', 'Occupation',
       'EmploymentStatus', 'EmploymentStatusDuration', 'IsBorrowerHomeowner',
       'CurrentlyInGroup', 'GroupKey', 'DateCreditPulled',
       'CreditScoreRangeLower', 'CreditScoreRangeUpper',
       'FirstRecordedCreditLine', 'CurrentCreditLines', 'OpenCreditLines',
       'TotalCreditLinespast7years', 'OpenRevolvingAccounts',
       'OpenRevolvingMonthlyPayment', 'InquiriesLast6Months', 'TotalInquiries',
       'CurrentDelinquencies', 'AmountDelinquent', 'DelinquenciesLast7Years',
       'PublicRecordsLast10Years', 'PublicRecordsLast12Months',
       'RevolvingCreditBalance', 'BankcardUtilization',
       'AvailableBankcardCredit', 'TotalTrades',
       'TradesNeverDelinquent (percentage)', 'TradesOpenedLast6Months',
       'DebtToIncomeRatio', 'IncomeRange', 'IncomeVerifiable',
       'StatedMonthlyIncome', 'LoanKey', 'LoanCurrentDaysDelinquent',
       'LoanMonthsSinceOrigination', 'LoanNumber', 'LoanOriginalAmount',
       'LoanOriginationDate', 'LoanOriginationQuarter', 'MemberKey',
       'MonthlyLoanPayment', 'LP_CustomerPayments',
       'LP_CustomerPrincipalPayments', 'LP_InterestandFees', 'LP_ServiceFees',
       'LP_CollectionFees', 'LP_GrossPrincipalLoss', 'LP_NetPrincipalLoss',
       'LP_NonPrincipalRecoverypayments', 'PercentFunded', 'Recommendations',
       'InvestmentFromFriendsCount', 'InvestmentFromFriendsAmount',
       'Investors'],
      dtype='object')
```

```
## List of identifiers and other irrelevant columns
```

```
identifiers = ["ListingKey", "ListingNumber", "GroupKey", "LoanKey", "LoanNumber", "MemberKey"]
df.drop(columns= identifiers, axis= 1, inplace= True) ## Dropping irrelevant columns
```

```
df.shape ##
```

```
→ (113066, 63)
```

The drop is done, since the features count went down to 63.

```
get_percent_null(df)
```

BorrowerAPR	0.022111
CreditScoreRangeUpper	0.522704
CreditScoreRangeLower	0.522704
InquiriesLast6Months	0.616454
TotalCreditLinespast7years	0.616454
FirstRecordedCreditLine	0.616454
PublicRecordsLast10Years	0.616454
CurrentDelinquencies	0.616454
DelinquenciesLast7Years	0.875595

```
TotalInquiries          1.025065
EmploymentStatus        1.994410
Occupation              3.121186
BorrowerState           4.877682
AvailableBankcardCredit 6.672209
TradesOpenedLast6Months 6.672209
TotalTrades             6.672209
TradesNeverDelinquent (percentage) 6.672209
PublicRecordsLast12Months 6.725276
RevolvingCreditBalance 6.725276
BankcardUtilization     6.725276
OpenCreditLines          6.725276
CurrentCreditLines       6.725276
AmountDelinquent         6.741195
EmploymentStatusDuration 6.743849
DebtToIncomeRatio        7.492969
ProsperScore             25.723029
ProsperRating (Alpha)    25.723029
ProsperRating (numeric)  25.723029
EstimatedReturn          25.723029
EstimatedLoss            25.723029
EstimatedEffectiveYield  25.723029
dtype: float64
```

Let's handle the lower percentages by only dropping the NA values instead of the whole columns.

```
df.dropna(inplace=True)
```

```
get_percent_null(df)
```

```
→ Series([], dtype: float64)
```

No missing NA values left.

```
df.shape
```

```
→ (75486, 63)
```

After handling missing values and dropping unnecessary columns the shape of the data is 76,216 rows and 63 columns

We want to create a simplified version of this columns, since there are multiple values for the Due Past values.

```
df['LoanStatus'].unique()
```

```
→ array(['Current', 'Past Due (1-15 days)', 'Defaulted', 'Completed',
       'Chargedoff', 'Past Due (16-30 days)', 'Past Due (61-90 days)',
       'Past Due (31-60 days)', 'Past Due (91-120 days)',
       'FinalPaymentInProgress', 'Past Due (>120 days)'], dtype=object)
```

```
df['SimplifiedLoanStatus'] = df['LoanStatus'].apply(lambda x: 'Past Due' if 'Past Due' in x
df.head(10)
```

	ListingCreationDate	Term	LoanStatus	BorrowerAPR	BorrowerRate	LenderYield	Estim
1	2014-02-27 08:28:07.900	36	Current	0.12016	0.0920	0.0820	
3	2012-10-22 11:02:35.010	36	Current	0.12528	0.0974	0.0874	
4	2013-09-14 18:38:39.097	36	Current	0.24614	0.2085	0.1985	
5	2013-12-14 08:26:37.093	60	Current	0.15425	0.1314	0.1214	
6	2013-04-12 09:52:56.147	36	Current	0.31032	0.2712	0.2612	
7	2013-05-05 06:49:27.493	36	Current	0.23939	0.2019	0.1919	
8	2013-12-02 10:43:39.117	36	Current	0.07620	0.0629	0.0529	
10	2012-05-10 07:04:01.577	60	Current	0.27462	0.2489	0.2389	
12	2013-12-15 20:01:10.757	36	Past Due (1-15 days)	0.17969	0.1435	0.1335	
13	2013-07-15 16:28:28.087	36	Current	0.13138	0.1034	0.0934	

10 rows × 64 columns

▼ Main Features of interest

- All the left columns are the features of interest they will be divided based on the objective into three subsets.
- Objective 1: Loan Performance

```
loan_performance_columns = [ 'LoanStatus','SimplifiedLoanStatus' , 'LoanCurrentDaysDelinquent',
 'LoanOriginalAmount', 'BorrowerAPR', 'DebtToIncomeRatio', 'CreditScoreRangeLower','CreditScoreRangeUpper',
 'ProsperScore', 'EmploymentStatus','IsBorrowerHomeowner' ]  
  
loan_performance_df = df[loan_performance_columns]  
print(loan_performance_df.shape)  
loan_performance_df.head()
```

→ (75486, 12)

	LoanStatus	SimplifiedLoanStatus	LoanCurrentDaysDelinquent	LoanMonthsSinceOriginating
1	Current	Current	0	
3	Current	Current	0	
4	Current	Current	0	
5	Current	Current	0	
6	Current	Current	0	

- Objective 2: Credit Score and Borrower Analysis

```
credit_borrower_columns = [ 'CreditScoreRangeLower', 'CreditScoreRangeUpper', 'ProsperRating',
 'ProsperRating (Alpha)', 'ProsperScore', 'IncomeRange', 'EmploymentStatus', 'IsBorrowerHomeowner',
 'BorrowerAPR', 'LoanOriginalAmount' ]  
  
credit_borrower_df = df[credit_borrower_columns]  
print(credit_borrower_df.shape)  
credit_borrower_df.head()
```

→ (75486, 10)

	CreditScoreRangeLower	CreditScoreRangeUpper	ProsperRating (numeric)	ProsperRating (Alpha)	ProsperScore
1	680.0	699.0	6.0	A	
3	800.0	819.0	6.0	A	
4	680.0	699.0	3.0	D	
5	740.0	759.0	5.0	B	
6	680.0	699.0	2.0	E	

- Objective 3: Geographic and Demographic Analysis

```
geo_demo_columns = [ 'BorrowerState', 'Occupation', 'EmploymentStatus', 'IncomeRange',
    'IsBorrowerHomeowner', 'LoanOriginalAmount', 'ProsperRating (Alpha)', 'CreditScoreRange',
    'CreditScoreRangeUpper' ]  
  
geo_demo_df = df[geo_demo_columns]  
print(geo_demo_df.shape)  
geo_demo_df.head()
```

→ (75486, 9)

	BorrowerState	Occupation	EmploymentStatus	IncomeRange	IsBorrowerHomeowner	LoanO
1	CO	Professional	Employed	\$50,000- 74,999		False
3	GA	Skilled Labor	Employed	\$25,000- 49,999		True
4	MN	Executive	Employed	\$100,000+		True
5	NM	Professional	Employed	\$100,000+		True
6	KS	Sales - Retail	Employed	\$25,000- 49,999		False

Store the cleaned dataset

```
df.to_csv('./data/prosperLoanDataCleaned.csv') ## Load the csv into pandas dataframe
```

▼ Univariate Exploration


In this section, we are investigating distributions of individual variables. To see unusual points or outliers, take a deeper look to clean things up and prepare yourself to look at relationships between variables. the "Question-Visualization-Observations" framework is used throughout the exploration, it involves asking a question from the data, creating a visualization to find answers, and then recording observations after each visualisation.

1. Loan Status

Question: What are the distribution of the Loan Statuses?

```
## add ordering to the loan status and include the cancelled ones
loan_status_order = ['Cancelled', 'Completed', 'FinalPaymentInProgress', 'Current',
                      'Past Due', 'Chargedoff', 'Defaulted']

plt.figure(figsize=(14, 6))
sns.countplot(data = loan_performance_df, y = 'SimplifiedLoanStatus', color = 'purple', order=loan_status_order)
plt.xlabel('Status')
plt.title('What are the distribution of the Loan Statuses?');
```



It seems that the majority status among the loan statuses is the current which is 50462, While the completed are success story they are 17675 . The chargedoff however are 4444 cases they are failed loans. While the deafulted are the cases in danger of chargedoff they reach 885. And the past due are slightly larger 1835 but still not in danger of charge off.


The chaged off and the deafulted represnts the actual risks for these loans.

Nothing unusual with this distribution.

2. The Loan Original Amount

Let's take a look into the Loan Original Amounts.

```
plt.figure(figsize=(14, 6))
sns.histplot(data= credit_borrower_df, x = 'LoanOriginalAmount', bins=20, kde= True, color
plt.title('What is the distribution of the Loan Original Amount?')
plt.ylabel('Count')
plt.xlim(1000)
plt.xlabel('Loan Original Amount');
```




The distribution of the Loan Original Amount is right skewed, with multiple peaks suggesting distinct borrower groups. Most loans are relatively small, but there are outliers indicating larger, less frequent loan amounts. However these are numerical outliers but valid values for a given loan amount, so doesn't need to be handled.

3. Borrower State

Which states are the borrowers from?

```
borrower_states_counts = geo_demo_df['BorrowerState'].value_counts()
```

```
plt.figure(figsize=(14, 6))
ax = sns.barplot(borrower_states_counts, color = 'purple')
plt.title('Which states are the borrowers from?')
plt.xlabel('State')
plt.xticks(rotation = 45);
```



California boasts the highest number of borrowers among all US states, followed closely by Texas, New York, and Florida. On the other end of the spectrum, states like South Dakota, Alaska, Vermont, and Wyoming have significantly lower borrower counts.

California is an unupper bound outlier, however this is numerically and valid values that will not be handled.

↙ Bivariate Exploration


In this section, relationships between pairs of variables in the data are investigated. The variables of intreset that have been introduced in the previous sections. Questions are asked.

Question: Does having investors affect the monthly loan payments?

```
df['Investors'].corr(df['MonthlyLoanPayment'])
```

→ 0.30761429000108165

```
plt.figure(figsize=(14, 6))
sns.scatterplot(data= df, x='MonthlyLoanPayment', y='Investors', alpha=0.3, color = 'purple'
plt.title('Does having investors affect the monthly loan payments?')
plt.xlabel('Monthly Loan Payment')
plt.ylabel('Investors');
```



There is a positive correlation between the number of investors and the monthly loan payments, with a correlation coefficient of 0.3. This indicates that as the number of investors increases, the monthly loan payments also tend to increase.


Some outliers are observed, with loans having over 600 investors and monthly payments exceeding 1000.

A few loans with very high monthly payments (above 2000) are also present, indicating significant loan amounts or terms.

The plot shows a dense cluster of points with fewer than 200 investors and monthly payments below 500, suggesting this is the most common scenario.

Question: Does Loan status have unusual loan original amount?

```
plt.figure(figsize=(14, 6))
sns.boxplot(data=loan_performance_df, x='SimplifiedLoanStatus', y='LoanOriginalAmount', colc
plt.title('Does Loan status have unusual loan original amount?')
plt.xlabel('Loan Status')
plt.ylabel('Loan Original Amount');
```



All loan statuses have outliers, indicating some loans are significantly larger than most others in their respective categories.


Completed loans exhibit the most outliers, suggesting a wide variation in loan amounts within this category.

Charged-off loans also show a considerable number of outliers.

Current Loans, The highest upper outliers are observed in the current loan status, indicating some very large loans are currently active.

Question: How does the Original Amount of Loan change over time?

```
plt.figure(figsize=(14, 6))
sns.lineplot(data= df, x = 'LoanOriginationDate', y ='LoanOriginalAmount', color = 'purple')
plt.ylabel('Amount')
plt.xlabel('Year')
plt.xlim(pd.to_datetime('2009-07-20'), pd.to_datetime('2014-03-12')) ## the max and min date
plt.title('How does the Original Loan Amount changes over time?');
```




There is a noticeable increase in loan amounts over the years, suggesting that borrowers have been taking out larger loans as time progresses.

Interestingly, the year 2009 stands out with some of the highest loan amounts, indicating that during this period, there were notably larger loans compared to other years.

Question: Does employment status have unusual credit score lower range?

```
employment_status_oredr = ['Employed', 'Full-time', 'Part-time', 'Self-employed', 'Retired',  
plt.figure(figsize=(14, 6))  
sns.boxplot(data=credit_borrower_df, x='EmploymentStatus', y='CreditScoreRangeLower', color=  
plt.title('Does employment status have unuaual Credit Score Range Lower?')  
plt.xlabel('Employment Status')  
plt.ylabel('Credit Score Range Lower');
```



Employed borrowers have a wide range of credit scores, with some higher outliers.

This indicates that while most employed borrowers have moderate to high credit scores, there are a few with exceptionally high scores.

Borrowers in the "Other" and "Full-time" categories have similar distributions of credit scores, generally higher than those of employed borrowers. These groups also include some high outliers, suggesting that a subset of these borrowers have excellent credit.

Self-employed borrowers tend to have lower credit scores compared to other groups, with a narrower range and fewer outliers.

Not employed borrowers have the lowest and least variable credit scores, indicating financial instability or limited credit history.

Question: What are the employment statuses of students who are taking loans?

```
## List of students occupations ordered by their natural order
study_occupations = ['Student - College Freshman', 'Student - College Sophomore', 'Student
'Student - Community College', 'Student - Technical School']


## Labels for the visualization
labels = ['College Freshman', 'College Sophomore', 'College Junior', 'Community College', '1

## the natural order of employment staus
employment_status_hue_oredr = ['Employed', 'Full-time', 'Part-time', 'Not employed']

plt.figure(figsize=(14, 6))
sns.countplot(data=geo_demo_df[geo_demo_df['Occupation'].isin(study_occupations)], x='Occupation',
               hue='EmploymentStatus', hue_order=employment_status_hue_oredr, order = study_occupations)
plt.title('What are the employment statuses of students who are taking loans?')
plt.xlabel('Students Occupation')
plt.ylabel('Count')
plt.xticks(ticks= study_occupations, labels= labels)
plt.legend(title='Employment Status');
```



What are the employment statuses of students who are taking loans?



Employment is widespread among students taking loans, with the majority falling into the "Employed" category across all educational levels.

Part-time work appears to be the most common arrangement, indicated by the higher count of "Part-time" compared to "Full-time" in most categories. Expected since they are students.

Community college students exhibit a notably higher proportion of full-time employment compared to other groups.


The "Technical School" category shows a lower overall employment rate and a higher percentage of students who are not employed.

Question: What are interesting relationships between variables in the dataset?


```
intrest_vars = ['LoanOriginalAmount', 'BorrowerAPR', 'BorrowerRate',
    'CreditScoreRangeLower', 'CreditScoreRangeUpper',
    'DebtToIncomeRatio', 'StatedMonthlyIncome']
subset = df[intrest_vars]
sns.heatmap(subset.corr(), annot=True)
plt.xticks(rotation = 45)
plt.title("Correlation Matrix");
```



Correlation Matrix



```
grid = sns.PairGrid(data = df, vars = intrest_vars )
grid.map_diag(plt.hist, bins = 20,color='purple')
grid.map_offdiag(plt.scatter,color='purple')
grid.tight_layout();
```



Positive correlation between borrowerAPR and BorrowerRate .

Debit Income Ratio has a low positive correlation with both the credit score range lower and the upper range .

LoanOriginalAmount and StatedMonthlyIncome have a moderate positive correlation (0.3). This indicates that individuals with higher incomes tend to borrow larger amounts.

BorrowerAPR and CreditScoreRangeLower have a moderate negative correlation (-0.55). This suggests that borrowers with higher credit scores tend to have lower APRs.

There is no other interesting correlations between the variables.


▼ Multivariate Exploration

In this section, plots of three or more variables are created to investigate the data even further.

Question: How does the original loan amount relate to the borrower's APR across different loan statuses?

```
## this list is the natural order of loan status  
loan_status_order = [ 'Completed', 'FinalPaymentInProgress', 'Current', 'Past Due', 'Chargedoff'
```

```
## Scatterplot with multiple encodings, color encoding with ordered categories
plt.figure(figsize=(14, 6))
sns.scatterplot(data=df, x='LoanOriginalAmount', y='BorrowerAPR', hue='SimplifiedLoanStatus'
                 hue_order=loan_status_order, palette='RdYlGn_r', s = 10)
plt.title('How does the original loan amount relate to the borrowers APR across different lc')
plt.ylabel('Borrower APR')
plt.xlabel('Loan Original Amount')
plt.legend(title = 'Loan Status', loc='upper right');
```




The majority of data points cluster in the lower left region, indicating that most loans have lower original amounts and APRs.

A few data points appear as outliers, situated away from the main cluster. These represent loans with significantly higher original amounts or APRs compared to the majority.

The current loans have different and varying APRs and loan amounts. While most of the completed, defaulted, and chargedoff are in the lower region.

Question: How does the original loan amount vary across different income ranges and loan statuses?

```
## Facet Plotting three variables SimplifiedLoanStatus, LoanOriginalAmount and IncomeRange
grid = sns.FacetGrid(data = df, col = 'SimplifiedLoanStatus', col_wrap = 3, col_order= loan_
grid.map(sns.barplot, 'LoanOriginalAmount', 'IncomeRange', color = 'Purple')
grid.set_titles("{col_name}")
grid.set_axis_labels('Loan Original Amount', 'Income Range')
plt.suptitle('How does the original loan amount vary across different income ranges and loan statuses?')
grid.figure.set_size_inches(14, 6);
```



As income range increases, the loan original amount also tends to increase. This is evident in the length of the bars across different income ranges.

The "Current" and "Completed" statuses, where the highest loan original amount is not always associated with the highest income range.

Higher income ranges might be associated with fewer defaults, which is worth investigating.

Across all loan statuses, the highest Loans amounts are associated with the 100,000+ income range.

Question: How do different terms group with the selected variables?

```
selected_columns = ['LoanOriginalAmount', 'BorrowerAPR', 'DebtToIncomeRatio']

## Plot Matrix
grid = sns.pairplot(data=df, vars=selected_columns, hue='Term', palette=['purple', 'yellow'],
plt.suptitle('How do different terms group with the selected variables?', y=1.02)
grid.figure.set_size_inches(14, 6);
```